

# Computer Vision

CVI620

Session 7

# Overview

---

Geometric Transformations

---

Affine

---

Scale, rotate, reflection, shear, translation

---

warpAffine

---

Noise Intro

# Agenda

---

Noise

---

Noise Types

---

Denoising Techniques

---

Filters

---

Convolution


---

Mean Denoising




# Noise

Definition: Anything that deviates from the ideal image or hinders achieving your imaging goal.



# Noise Source



# Noise Characteristics

- Additive and random
- Represented as:

$$P(i, j) = I(i, j) + n(i, j)$$

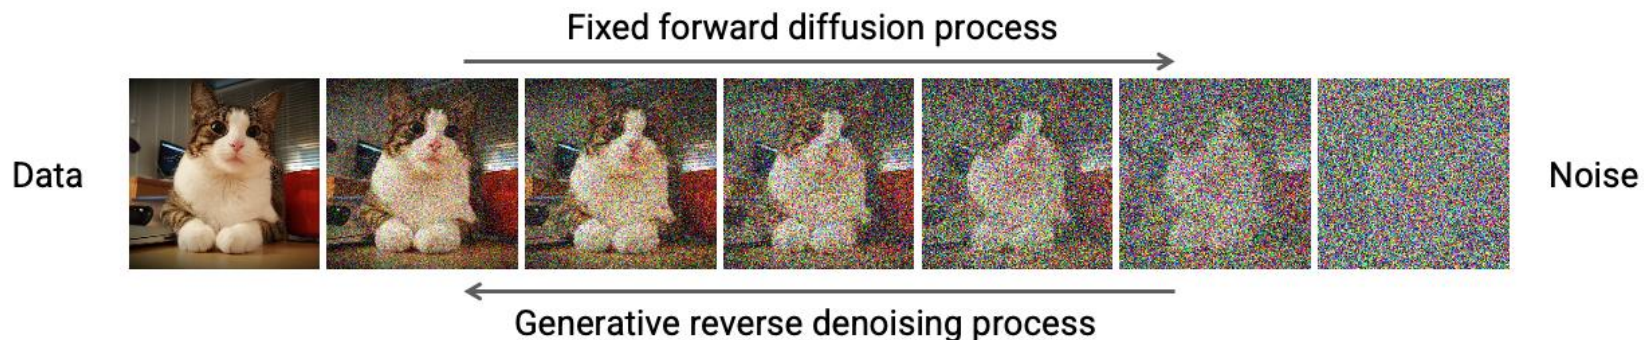
$P(i, j)$ : Pixel value in the noisy image

$I(i, j)$ : Pixel value in the ideal image

$n(i, j)$ : Noise value


# Why we need to study noise?

- Ensures image accuracy for analysis or display.
- Critical for applications in medical imaging, machine vision, and remote sensing.
- Used in state-of-art generative models like DALL-E



# Noise Types

- Gaussian noise
- Salt and Pepper
- Poisson noise
- Speckle noise
- Thermal noise






# Gaussian Noise

- Noise that follows a normal (Gaussian) distribution
- Normal distribution and why it is important:

- Symmetry: Centered around the mean ( $\mu$ ).
- Mean, Median, and Mode: All are equal and located at the peak.
- Width determined by standard deviation ( $\sigma$ ).



# Gaussian Noise in CV2

- Choose random samples from normal distribution
- Add it to the image

```
image = cv2.imread('Lucy.jpg')
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

mean = 0
std_dev = 25

gaussian_noise = np.random.normal(mean, std_dev, image.shape).astype('float32')

noisy_image = cv2.add(image.astype('float32'), gaussian_noise)

noisy_image = np.clip(noisy_image, 0, 255).astype('uint8')

plt.figure(figsize= (10, 5))
plt.subplot(1, 2, 1)
plt.imshow(image)

plt.subplot(1, 2, 2)
plt.imshow(noisy_image)

plt.show()
```

# Impulsive Salt and Pepper Noise

- A type of impulse noise where random pixels in the image are replaced with:
  - White (salt): Maximum intensity (e.g. 255 in an 8-bit image)
  - Black (pepper): Minimum intensity (e.g. 0 in an 8-bit image)
- Appears as random white and black dots in an image

$$P(i, j) = \begin{cases} I(i, j), & \text{with probability } (1 - p) \\ 0 \text{ (pepper)}, & \text{with probability } p/2 \\ 255 \text{ (salt)}, & \text{with probability } p/2 \end{cases}$$

# Salt and Pepper

```
import cv2
import numpy as np

image = cv2.imread('Lucy.jpg', cv2.IMREAD_GRAYSCALE)

def add_noise(img):
    row , col = img.shape

    number_of_pixels = np.random.randint(300, 10000)
    for i in range(number_of_pixels):
        y_coord = np.random.randint(0, row - 1)
        x_coord = np.random.randint(0, col - 1)
        img[y_coord, x_coord] = 255

    number_of_pixels = np.random.randint(300, 10000)
    for i in range(number_of_pixels):
        y_coord = np.random.randint(0, row - 1)
        x_coord = np.random.randint(0, col - 1)
        img[y_coord, x_coord] = 0

    return img

noisy_img = add_noise(image)
cv2.imshow('Noisy Image', noisy_img)
cv2.waitKey(0)
cv2.destroyAllWindows()
```


Question?

We have added noise so far.

How can we remove it?

# Denoising


- [https://www.topazlabs.com/denoise-ai?srsltid=AfmBOopM02\\_xy6QWHcr\\_WM5I5iSE0MH4MBimXXkvw7fg80zWAmOHNnUs](https://www.topazlabs.com/denoise-ai?srsltid=AfmBOopM02_xy6QWHcr_WM5I5iSE0MH4MBimXXkvw7fg80zWAmOHNnUs)




# Denoising Techniques



Spatial Filtering



Frequency Domain  
Filtering



Advanced Methods


# Spatial Filtering

- Mean Filter: Averages pixel values in a neighborhood
- Median Filter: Replaces pixel value with the median in a local window
- Gaussian Filter: Weighted averaging with a Gaussian kernel



# Frequency Domain Filtering

- Fourier Transform-Based Filtering: Suppresses high-frequency noise



# Advanced Methods

01

Wavelet  
Transform  
Denoising:  
Removes noise  
at different scales

02

Non-Local Means  
(NLM): Uses  
similar patches  
across the image

03

Deep Learning-  
Based Denoising  
(Denoising  
Autoencoders,  
CNNs)

# Another Categorization



1. Filtering Based Denoising



2. Transform Based Denoising



3. Statistical & Probabilistic Methods



4. Machine Learning & Deep Learning Approaches




5. Optimization-Based Methods


# Filtering Based Denoising

- Mean Filter
- Gaussian Filter
- Median Filter
- Bilateral Filter
- Wavelet Denoising

# Transformer Based Denoising




Fourier Transform Filtering:  
Removes noise in the frequency domain.



Wavelet Transform: Decomposes  
the image into different scales  
and removes noise selectively.

# Statistical Denoising



Non-Local Means (NLM):  
Uses patches from the  
image to reduce noise.



Bayesian Denoising:  
Uses probabilistic  
models to infer the  
denoised image.

# ML and DL Denoising

- Denoising Autoencoders: Neural networks trained to remove noise from images.
- Convolutional Neural Networks (CNNs): Trained models that learn patterns to reconstruct noise-free images.
- Generative Models (GANs, Diffusion Models): Learn distributions of clean images to remove noise effectively.

# Optimization Based Denoising

- Total Variation (TV) Denoising: Minimizes variations while preserving edges.
- Sparse Coding: Represents images in a sparse domain and removes noise adaptively.




noise -> filter -> convolution -> denoise

# Image Filtering

# Filter

- The core idea is to manipulate image features like noise, edges, and textures based on specific objectives.
- Adding, changing, detecting features or filters in a picture is better to be applied step by step and region by region




# Convolution

Filtering is commonly implemented using **convolution**, especially in spatial domain filtering.


In frequency domain filtering, convolution is performed using Fourier Transform methods.

# Convolution in Signal and Systems




# Convolution

- Sum of pixel multiplications in a region



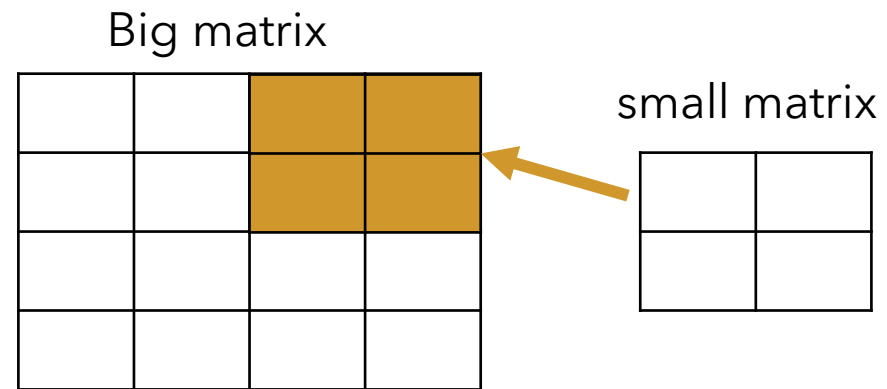
# Convolution

- Align from (0,0)



# Convolution

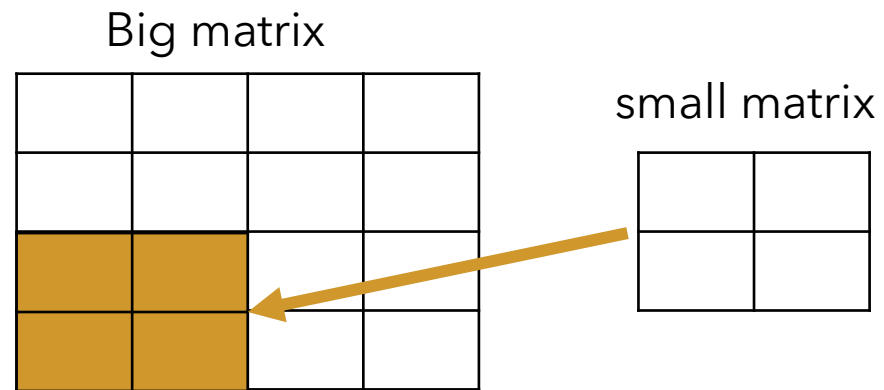
- Move left to right with a step size






# Convolution

- Move top to bottom with a step size




# Convolution




# Convolution

- Mathematical operation



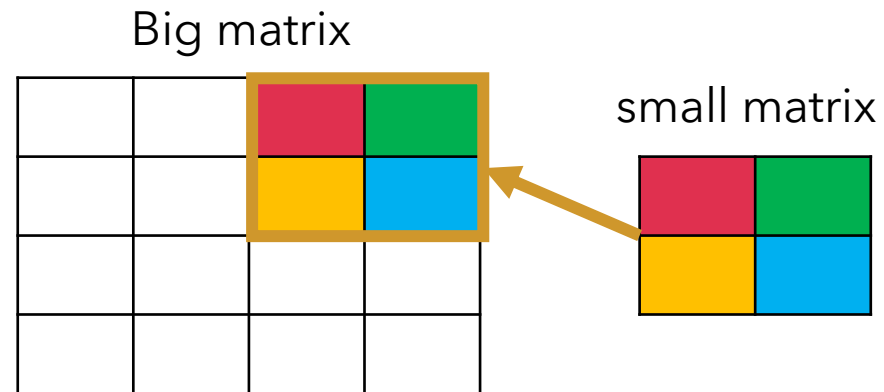
# Convolution

- Sum of pixel multiplications in a region
- Multiply aligned pixels and add them together



# Convolution

- Sum of pixel multiplications in a region
- Multiply aligned pixels and add them together



# Convolution

- Sum of pixel multiplications in a region
- Multiply aligned pixels and add them together
- **Make a new image**

Big matrix


small matrix


Final Image


# Example

Image

0	1	2
3	4	5
6	7	8

Kernel

0	1
2	3



19	25
37	43


$$0*0 + 1*1 + 3*2 + 3*4 = 19$$

$$1*0 + 2*1 + 4*2 + 5*3 = 25$$

...

# Real Example

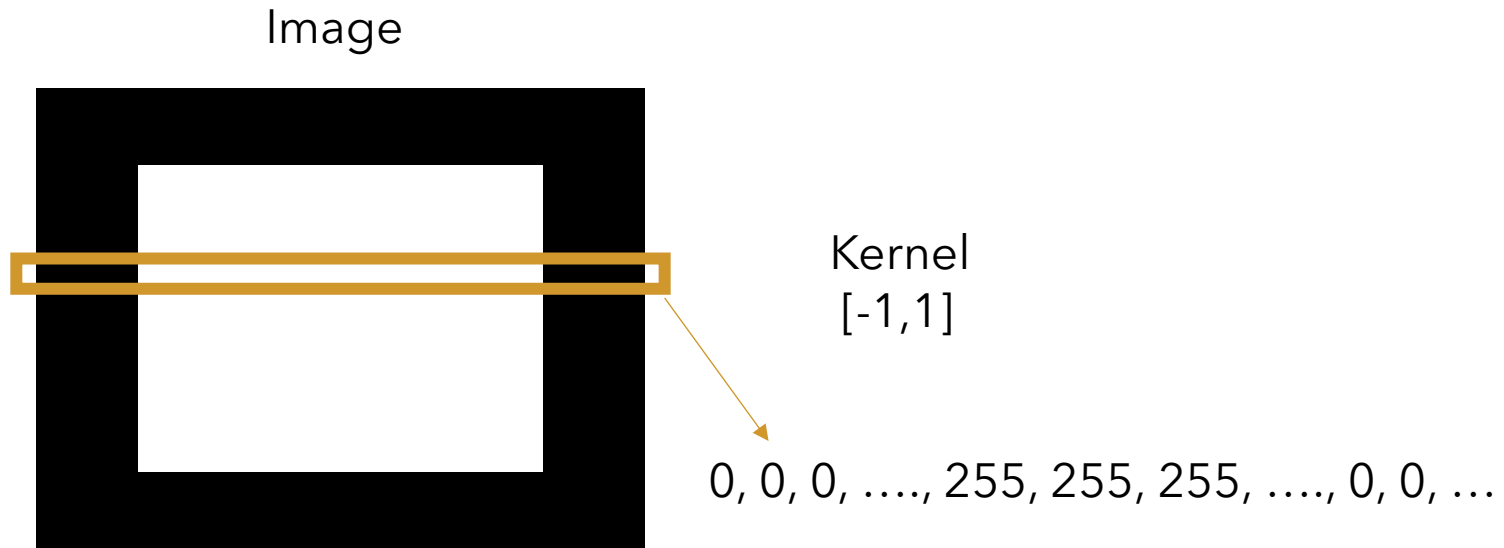
Image




Kernel  
[-1,1]



# Real Example



# Real Example



Kernel  
[-1,1]

0, 0, 0, ..., 0, 255, 255, 255, ..., 0, 0, ...

$$0 * -1 + 0 * 1 = 0$$

$$0 * 1 + 255 * 1 = 255$$


$$255 * -1 + 255 * 1 = 0$$

$$255 * -1 + 255 * 1 = 0$$

$$255 * -1 + 255 * 1 = 0$$

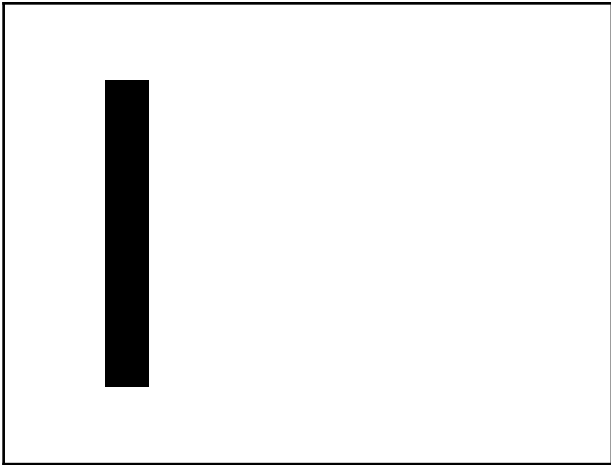
$$255 * -1 + 0 * -1 = -255 \rightarrow \text{clip}(-255) = 0$$

# Real Example



# Real Example

0, 0, 0, ..., 0, 255, 0, 0, 0, ...



Edge!!

# Let's Code

```
import cv2
import numpy as np

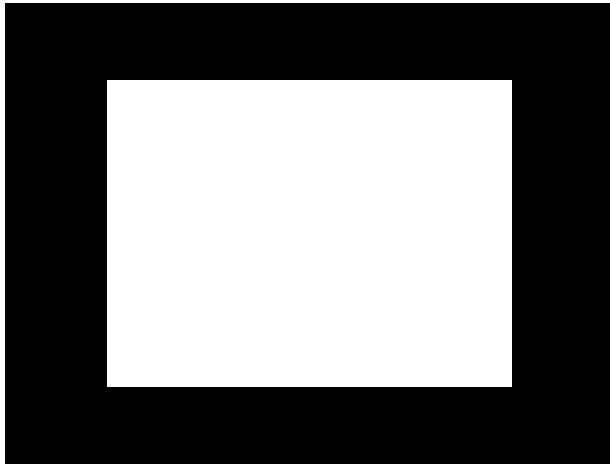
img = cv2.imread("square.png")
kernel = np.array([[-1, 1]])

out_image = cv2.filter2D(img, cv2.CV_8U, kernel)


cv2.imshow("left edge", out_image)
cv2.waitKey(0)
```

# Let's Convolve Vertically

Image



Kernel  
[[-1],  
[1]]



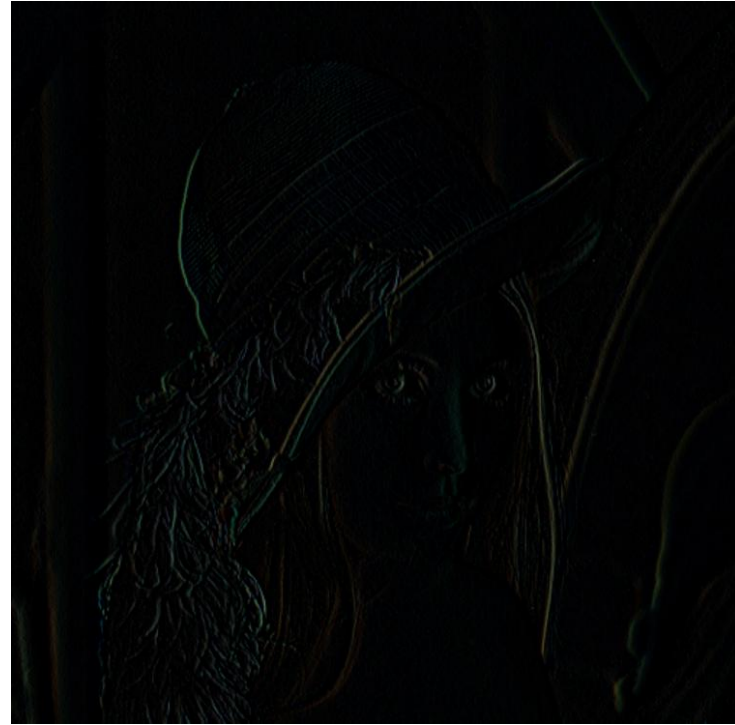
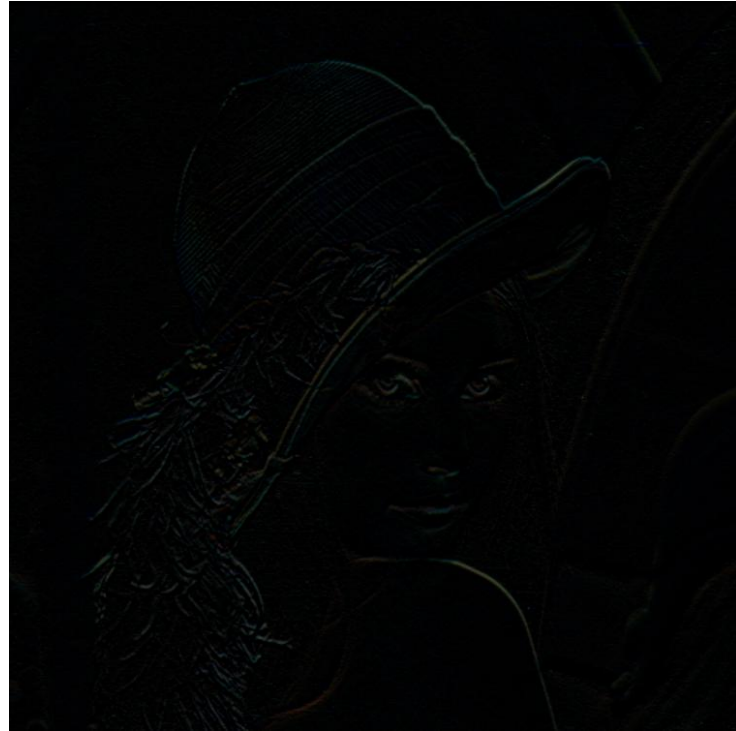
# Example

```
import cv2
import numpy as np

img = cv2.imread("Lenna.png")
kernel1 = np.array([[ -1, 1]])
kernel2 = np.array([[ -1],
                    [ 1]])

out_image1 = cv2.filter2D(img, cv2.CV_8U, kernel1)
out_image2 = cv2.filter2D(img, cv2.CV_8U, kernel2)

cv2.imshow("left edge", out_image1)
cv2.imshow("bottom edge", out_image2)
cv2.waitKey(0)
cv2.destroyAllWindows()
```



# More Complex Convolutions!



-1	0	1
-2	0	2
-1	0	1

-1	-2	-1
0	0	0
1	2	1






# How to achieve these kernels?



Experiments!



ML to learn kernels!

Now let's get back to denoising!

# Mean Denoising


- Replaces each pixel value with the average of its neighboring pixels.
- Smooths the image by averaging pixel intensities.
- Reduces random noise while preserving structure.



# Mean Denoising



The filter slides over the image, replacing each pixel with the mean of its neighbors.



Uses a convolution operation with a kernel

# Mean Denoising

- Applying a 3×3 Mean Filter
- Kernel Example:

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} * 1/9$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Each pixel value is replaced by the average of surrounding pixels.
- Result: A smoother image with reduced noise.
- Also called blurring!

# Example

```
import cv2
import numpy as np

image = cv2.imread("Lucy.jpg")

if image is None:
    print("ERROR! Image not available...!")

# Gaussian Noise
noise = np.random.normal(0, 25, image.shape).astype('float32')

noisy_image = image + noise
noisy_image = np.clip(noisy_image, 0, 255).astype('uint8')

kernel = np.ones((3,3), dtype= np.float32) / 9
denoised_image = cv2.filter2D(noisy_image, -1, kernel)

cv2.imshow("left edge", denoised_image)
cv2.imshow('frame', noisy_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

